

3we: An Open Infrastructure for Sim-to-Real Embodied AI Research

3we Contributors

Open Source Project

<https://github.com/telleroutlook/3we-robot-platform>

Abstract—Embodied AI research requires tight integration of simulation, real hardware, and learning algorithms—yet existing platforms force researchers to choose between expensive proprietary systems, simulation-only environments, or hardware platforms with steep ROS2 learning curves. We present 3we, a fully open-source infrastructure that exposes a consistent Python API across a lightweight mock simulator, Gazebo, NVIDIA Isaac Sim, and real hardware (Raspberry Pi 5 + Hailo-8L), so the same user code targets every backend without modification. We make no claim of zero-effort sim-to-real transfer: closing the residual gap (sensor models, dynamics, actuator delay) still requires task-specific tuning, which the platform supports through documented domain-randomization and fault-injection hooks. The platform provides an AI-First Python API that reduces typical ROS2 navigation code from 60+ lines to 5; four Gymnasium-compatible environments (Navigation, Exploration, ObjectNav, Vision-Language Navigation) plus a multi-agent cooperative environment; native VLM/VLA integration with Hailo NPU edge inference; a Hardware Abstraction Layer (HAL) supporting third-party robots; trajectory recording with LeRobot/HuggingFace Hub interoperability; and standardized benchmarks across 7 evaluation scenes. On reproducible hardware costing approximately \$300 (BOM), preliminary validation yields Sim-to-Real transfer ratios exceeding 0.6 on point navigation tasks, with community baselines achieving 82% success rate (SPL 0.65) in structured simulation environments. All software (Apache 2.0), hardware designs (CERN-OHL-P v2), and documentation (CC-BY-SA 4.0) are publicly available.

I. Introduction

Embodied AI—the study of intelligent agents that perceive and act in physical environments—has emerged as a central challenge in artificial intelligence [1], [2]. Recent advances in Vision-Language Models (VLMs) [3] and Vision-Language-Action (VLA) models [4], [5] have demonstrated that foundation models can directly control robots through natural language instructions. However, translating these advances into reproducible research faces three fundamental barriers:

(1) **Hardware cost and accessibility.** Leading mobile robot platforms (TurtleBot 4: \$1,200+; Stretch RE1: \$25,000) price out most research groups, particularly in developing countries and teaching contexts.

(2) **Simulation-to-Reality gap.** Researchers typically develop in simulation (Habitat [2], Isaac Sim [6]) but face significant engineering effort to deploy on real hardware, often requiring complete code rewrites.

(3) **Software complexity.** ROS2 [7] provides robust middleware but imposes a steep learning curve on AI




Fig. 1. Code complexity comparison: Raw ROS2 vs. 3we Python API for common tasks. The API achieves 10–12× reduction in lines of code.

researchers who want to focus on algorithms rather than communication infrastructure.

We introduce 3we (“three-way Embodied”), an open infrastructure that addresses all three barriers simultaneously:

- **Open hardware** at approximately \$300 BOM with fully reproducible PCB designs (KiCad), mechanical drawings, and sourcing guides under CERN-OHL-P v2.
- **Consistent API** across backends: a unified Python API where `Robot(backend="gazebo")` and `Robot(backend="real")` execute the same user code without modification. Closing the residual sim-to-real gap remains a task-level concern, supported by the platform’s noise and fault-injection hooks.
- **AI-First API**: 5 lines from installation to first navigation, with native VLM/VLA integration, multi-level action spaces, and Gymnasium compatibility.
- **Standardized benchmarks**: 3 task types across 7 scenes with reproducible baselines and community leaderboard.
- **Hardware Abstraction Layer**: built-in profiles for 3we, AgileX Scout, TurtleBot4, Unitree Go2, with user-extensible YAML definitions.
- **Data and reproducibility**: trajectory recorder exports to LeRobot/Open-X format; ExperimentProtocol locks seed, version, and hyperparameters.

TABLE I
Platform Comparison

Feature	3we	TurtleBot4	Isaac Lab	LeRobot	Habitat
Open Hardware	✓	–	–	✓	–
BOM ≤\$300	✓	–	N/A	✓	N/A
Cross-backend API	✓	–	Sim only	–	Sim or
ROS2 + Nav2	✓	✓	Partial	–	–
AI-First Python	✓	–	✓	✓	✓
VLM/VLA Native	✓	–	–	✓	–
Edge AI (NPU)	✓	–	N/A	–	N/A
Gymnasium Env	✓	–	✓	✓	✓

II. Related Work

A. Mobile Robot Platforms

TurtleBot [8] has been the de facto standard for ROS education since 2010. TurtleBot 4 (\$1,200+) offers Nav2 integration but lacks AI-first APIs, simulation-hardware consistency guarantees, and costs 4× our platform. Linorobot2 [9] provides open designs but targets hobbyists without research-grade benchmarks.

B. Simulation Environments

Habitat [2] and AI2-THOR [10] provide high-fidelity visual navigation but lack physical robot deployment paths. NVIDIA Isaac Sim [6] enables GPU-accelerated RL training but requires expensive compute and offers no open hardware reference. These systems address simulation or reality—not the bridge between them.

C. AI Robot Frameworks

LeRobot [11] (HuggingFace) provides VLA model deployment for manipulation but lacks ROS2 integration, autonomous navigation stacks, and Sim2Real guarantees for mobile platforms. Its recent mobile extensions (LeKiwi, EarthRover) remain early-stage without Nav2-level autonomy.

D. Positioning

3we occupies a unique position: fully open hardware and software with a guaranteed-consistent Python API across all backends (an API-level contract; sim-to-real transfer still requires task-specific tuning), ROS2-native navigation, and AI-first design. Table I summarizes key differences.

III. System Architecture

A. Three-Layer Design

3we adopts a layered architecture that decouples AI research from robotics engineering:

- 1) AI-First Python API (user-facing): Researchers write standard Python with `async/await`. No ROS2, launch files, or message types exposed.
- 2) 3we-core (middleware): Manages backend dispatch, sensor data standardization, coordinate transforms, safety checks, hardware abstraction, and domain randomization noise injection.




Fig. 2. Three-layer architecture. Users write Python against the AI-First API (top). The middleware dispatches to one of four backends (bottom) with identical semantics. ROS2 infrastructure is fully transparent.

- 3) ROS2 / micro-ROS (infrastructure): Nav2 path planning, SLAM, micro-ROS bridge to ESP32 firmware. Entirely transparent to users.

B. Backend Polymorphism

The core abstraction is BackendBase, implemented by four backends:

- Mock: Zero-dependency 2D kinematic simulation with configurable noise and collision detection. Runs anywhere, no GPU or ROS2 needed.
- Gazebo: Full physics via Gazebo Harmonic. CPU-friendly, CI-compatible (headless mode).
- Isaac Sim: GPU-accelerated rendering and physics for large-scale parallel RL training with domain randomization and vectorized policy optimization.
- Real: Physical hardware via ROS2 topics. Identical API surface.

All backends return typed dataclasses (Pose2D, LaserScan, RGBDImage, IMUData, BatteryState, etc.) with identical fields and semantics, enforced by shared type definitions.

C. API Design

The following example demonstrates the complete workflow for VLM-powered navigation:

```
from threewe import Robot

async with Robot(backend="gazebo") as robot:
    image = robot.get_camera_image() # (H,W,3) uint8
    scan = robot.get_lidar_scan()    # LaserScan
    result = await robot.execute_instruction(
        "find the red door and move toward it"
    )
    print(f"Success: {result.success}")
```

Switching to real hardware requires only changing the backend parameter—no code modification, no recompilation, no configuration changes.

The API provides a rich perception and action interface: synchronous perception (`get_camera_image`,




Fig. 3. Mock backend visualization: autonomous navigation in the office_v2 scene. The robot (blue) follows a planned path (orange) while 360° LiDAR rays (green) detect obstacles. This zero-dependency simulation runs on any machine without GPU or ROS2.

TABLE II
Standard SKU Hardware (BOM: approximately \$300)

Component	Selection	Rationale
Compute	Raspberry Pi 5 (8GB)	Best price/performance ARM SBC
AI Accelerator	Hailo-8L (13 TOPS)	Edge inference for VLM/VLA
MCU	ESP32-S3 + micro-ROS	Real-time control + ROS2 native
Ranging	4× HC-SR04 ultrasonic	Basic obstacle avoidance (LiDAR optional)
IMU	BNO055 (9-axis)	On-chip fusion, zero configuration
Drive	4× 65mm Mecanum wheels	Omnidirectional, indoor-optimized
Camera	1080p USB fisheye 170°	Wide FoV for VLM tasks
Safety	Dual-channel relay + E-stop	ISO 13850, hardware-enforced cutoff
Power	3× 2S 18650 + BMS	7.4V supply, extended runtime

get_lidar_scan, get_pose, get_imu, get_battery_state, get_map, get_wheel_speeds, get_motor_current), async navigation (move_to, move_forward, rotate, explore, follow_path), and a standardized observation interface for RL (get_observation(modalities=...)). VLA policy outputs drive the robot directly via execute_action(action).

D. Hardware Specification

Table II summarizes the reference hardware at standard SKU level.

Optional upgrade: LD06 LiDAR (+\$17) for 2D SLAM and Nav2 autonomous navigation.

All designs are released under CERN-OHL-P v2 with KiCad 8 source files, Gerber manufacturing outputs, and step-by-step assembly documentation.

IV. Sim-to-Real Consistency

A. Sensor Noise Model

The mock and Gazebo backends inject physically-grounded noise matching the reference hardware:

- LiDAR (LD06): Gaussian range noise $\sigma = 8$ mm, 1% missing ray probability




Fig. 4. Sim-to-Real transfer ratios across five standard validation tests. All tests pass the minimum threshold (dashed line). Higher is better (1.0 = perfect transfer).

TABLE III
Sim2Real Transfer Tests

Test	Metric	Pass Criterion	Ratio
Straight walk 5m	Endpoint error	$\text{real} \leq \text{sim} \times 2.0$	0.85
360° rotation	Angle error	$\text{real} < 1.0^\circ$	0.92
Obstacle avoid (3)	Success rate	$\text{real} \geq \text{sim} \times 0.7$	0.78
PointNav 10m	SPL	$\text{real} \geq \text{sim} \times 0.6$	0.72
Dynamic obstacle	Collision rate	$\text{real} \leq \text{sim} \times 1.5$	0.81

- IMU (BNO055): Gyroscope noise $\sigma = 0.0014$ rad/s, accelerometer noise $\sigma = 0.01$ m/s²
- Camera (OV5647): Pixel noise $\sigma = 3$, brightness drift $\pm 5\%$
- Odometry: Mecanum wheel slip 5–15% stochastic variation, heading noise $\sigma = 0.002$ rad
- Motor current: Measurement noise $\sigma = 15$ mA, fixed offset bias 5 mA

B. Transfer Validation Protocol

We define five standard transfer tests with explicit pass criteria:

Overall pass rate requires $\geq 80\%$ of tests passing. The transfer ratios in Table III are derived from preliminary hardware trials with the documented noise model; large-scale validation is ongoing.

Transfer validation can be executed via CLI with a single command:

```
threewe sim2real report --backend gazebo \
--scene office_v2 --trials 5
```

V. Benchmark Suite

A. Tasks

The benchmark suite defines three navigation tasks of increasing difficulty:

- PointNav: Navigate to (x, y) coordinates (success threshold: 0.5m). Metrics: Success Rate (SR), SPL, average duration. Valid across all 7 scenes.
- ObjectNav: Find a target object from 10 categories (chair, desk, door, plant, monitor, etc.). Metrics: SR, SPL. Valid in 5 scenes.




Fig. 5. PointNav baseline results (Nav2 DWB planner, 100 episodes per scene). Success Rate and SPL decrease with scene complexity.

TABLE IV
Nav2 Baseline Results (100 episodes per configuration)

Task	Scene	SR $\uparrow$	SPL $\uparrow$	Duration (s)
PointNav	office_v2	0.82	0.65	12.3
PointNav	apartment_v1	0.68	0.52	15.7
PointNav	corridor_v1	0.91	0.78	18.5
PointNav	warehouse_v1	0.75	0.58	20.1
PointNav	cluttered_v1	0.60	0.42	14.8
PointNav	dynamic_v1	0.52	0.35	18.2
ObjectNav	office_v2	0.55	0.38	22.5
ObjectNav	apartment_v1	0.42	0.30	28.3
Exploration	office_v2	—	—	95.0
Exploration	warehouse_v1	—	—	110.0

Exploration metric: Coverage (office: 0.85, warehouse: 0.78)

- Exploration: Maximize area coverage within time budget (threshold: 80%). Metrics: Coverage fraction, duration. Valid in 5 scenes.

B. Evaluation Environments

Seven standardized scenes span diverse settings: office_v2 (20×15m), apartment_v1 (8×10m), corridor_v1 (50m linear), warehouse_v1 (30×20m), cluttered_v1 (5×5m), outdoor_v1 (50×30m), and dynamic_v1 (15×15m, moving obstacles).

Each scene provides 10–50 predefined start/goal pairs for reproducibility, with seeds ranging [0, 99] for 100-episode evaluation runs.

C. Baseline Results

Table IV presents Nav2 DWB planner baselines across representative scenes.

These baselines serve as reference points for the public community leaderboard, which is hosted as a versioned JSON file in the repository (data/leaderboard.json) and rendered at <https://3we.org/benchmarks>. Submissions are validated against the standardized seed protocol and SDK version. Results are fully reproducible via:

```
threewe benchmark run --task pointnav \
  --scene office_v2 --episodes 100
```

D. Gymnasium Integration

For RL researchers, 3we provides four standard Gymnasium environments and a multi-agent cooperative environment:

```
import gymnasium as gym

# Point-to-point navigation
env = gym.make("3we/Navigation-v1")

# Coverage exploration
env = gym.make("3we/Exploration-v1")

# Object-goal navigation (10 categories)
env = gym.make("3we/ObjectNav-v1")

# Vision-Language Navigation (instruction-following)
env = gym.make("3we/VLN-v1")

obs, info = env.reset()
obs, reward, term, trunc, info = env.step(action)
```

The Navigation environment supports three action levels (CMD_VEL: end-to-end RL; WHEEL_VEL: low-level control research; WAYPOINT: high-level VLM/LLM planning) and returns low-level physics data (wheel speeds, IMU, motor current) in the info dict for domain randomization research.

Additionally, MultiAgentEnv implements a PettingZoo-compatible parallel API for 2–4 robot cooperative exploration training.

VI. Hardware Abstraction and Data Infrastructure

A. Hardware Abstraction Layer (HAL)

To support heterogeneous robot platforms, 3we implements a HardwareProfile protocol that decouples kinematic parameters (wheel type, separation, radius, velocity limits) and sensor configuration from control logic. Built-in profiles cover four platforms:

- 3we_standard_v2: Mecanum wheels, 0.24m separation, max 0.5 m/s
- AgileX Scout: Differential drive, 65 kg, max 1.5 m/s
- TurtleBot4: Differential drive, 3.9 kg, max 0.31 m/s
- Unitree Go2: Quadruped, 15 kg, max 2.5 m/s

Researchers can define custom hardware profiles via YAML files, enabling Robot(hardware="my_robot") plug-and-play integration.

B. Trajectory Recording and Data Sharing

The threewe.data module provides:

- TrajectoryRecorder: Records observation-action pairs at fixed frequency during teleoperation or autonomous execution.
- HDF5/LeRobot export: Dataset format compatible with HuggingFace Hub, directly usable for VLA model training.
- Dataset card generation: Automatic README.md metadata conforming to Hub specifications.

C. Domain Randomization and Fault Injection

The `threewe.sim` module provides hardware-calibrated sensor noise models (`SensorNoiseModel`) covering LiDAR range noise, IMU drift, camera pixel noise, odometry slip, and motor current perturbation. Users inject noise in training loops via `noise.apply_lidar(scan)` without modifying the underlying simulator:

```
from threewe.sim.noise import SensorNoiseModel
noise = SensorNoiseModel(seed=42)
noisy_scan = noise.apply_lidar(clean_scan)
noisy_accel, noisy_gyro, noisy_ori = noise.apply_imu(
    acceleration, angular_velocity, orientation)
```

Beyond sensor-level noise, `FaultInjector` wraps any backend to simulate communication-layer faults—configurable latency with jitter, stochastic packet loss (returning stale sensor data), and timed disconnections—enabling chaos-engineering tests of policies and ROS2 stacks without modifying user code:

```
from threewe.sim.fault_injection import (
    FaultConfig, FaultInjector)
config = FaultConfig(latency_ms=50,
    packet_loss_rate=0.1)
faulty = FaultInjector(backend, config)
```

D. Experiment Reproducibility

The `ExperimentProtocol` dataclass locks all experiment variables—random seed, backend, scene, hardware profile, task, and hyperparameters—and automatically records SDK version and Git SHA, ensuring exact reproducibility.

VII. Conclusion and Future Work

We presented 3we, a fully open infrastructure for Embodied AI research that bridges the gap between simulation and reality through a unified Python API. Key contributions include: (1) a Python API that runs unchanged across four backends (mock, Gazebo, Isaac Sim, real hardware), with explicit hooks for closing the residual sim-to-real gap rather than a claim of zero-effort transfer; (2) reproducible open hardware at approximately \$300 BOM; (3) AI-first API design reducing navigation code complexity by 12×; (4) standardized benchmarks with community baselines and a public leaderboard; (5) Hardware Abstraction Layer supporting third-party robots; (6) trajectory recording with LeRobot/Open-X format interoperability; and (7) a domain-randomization module covering both sensor noise and communication-layer fault injection, with Hailo NPU edge VLA inference support.

The project ships 6 Jupyter Notebook tutorials (from Hello World to sim-to-real odometry drift analysis), enabling researchers to immediately verify cross-backend API consistency without hardware.

Future directions include: large-scale vectorized RL training using Isaac Sim parallel environments, cross-robot zero-shot policy transfer (leveraging HAL to normalize heterogeneous kinematics), deeper integration with

RoboCup standard tasks, and end-to-end foundation model fine-tuning for open-world navigation.

The platform is available at <https://github.com/telleroutlook/3we-robot-platform> under Apache 2.0 (software), CERN-OHL-P v2 (hardware), and CC-BY-SA 4.0 (documentation).

References

- [1] P. Anderson, A. Chang, D. S. Chaplot, A. Dosovitskiy, S. Gupta, V. Koltun, J. Kosecka, J. Malik, R. Mottaghi, M. Savva et al., “On evaluation of embodied navigation agents,” in arXiv preprint arXiv:1807.06757, 2018.
- [2] M. Savva, A. Kadian, O. Maksymets, Y. Zhao, E. Wijmans, B. Jain, J. Straub, J. Liu, V. Koltun, J. Malik et al., “Habitat: A platform for embodied ai research,” in Proceedings of the IEEE/CVF International Conference on Computer Vision, 2019, pp. 9339–9347.
- [3] OpenAI, “Gpt-4v(ision) system card,” OpenAI Technical Report, 2023.
- [4] A. Brohan, N. Brown, J. Carbajal, Y. Chebotar, X. Chen, K. Choromanski, T. Ding, D. Driess, A. Dubey, C. Finn et al., “Rt-2: Vision-language-action models transfer web knowledge to robotic control,” arXiv preprint arXiv:2307.15818, 2023.
- [5] K. Black, N. Brown, D. Driess, A. Esber, M. Suber, B. Ichter, P. Sermanet et al., “ π_0 : A vision-language-action flow model for general robot control,” arXiv preprint arXiv:2410.24164, 2024.
- [6] NVIDIA Corporation, “Isaac sim: Scalable robotics simulation,” NVIDIA Developer, 2023.
- [7] S. Macenski, T. Foote, B. Gerkey, C. Lalancette, and W. Woodall, “Robot operating system 2: Design, architecture, and uses in the wild,” *Science Robotics*, vol. 7, no. 66, p. eabm6074, 2022.
- [8] Open Robotics, “Turtlebot,” <https://www.turtlebot.com/>, 2023.
- [9] Linorobot, “Linorobot2: Autonomous mobile robot framework,” <https://linorobot.org/>, 2023.
- [10] E. Kolve, R. Mottaghi, W. Han, E. VanderBilt, L. Weihs, A. Herrasti, D. Gordon, Y. Zhu, A. Gupta, and A. Farhadi, “Ai2-thor: An interactive 3d environment for visual ai,” in arXiv preprint arXiv:1712.05474, 2017.
- [11] R. Cadene, S. Alibert, P. Sermanet et al., “Lerobot: State-of-the-art machine learning for real-world robotics,” GitHub repository, 2024.